


Redes Neurais Convolucionais



Input Image
224 x 224 x 3



Number of Trainable Weights

$$W_1 = 150,528 \times 128 = 19,267,584 \text{ Weights}$$

$$W_2 = 19,267,584 \times 128 \sim 2.4 \text{ Billion}$$

$$W_3 = 2.4B \times 128 > 300 \text{ Billion}$$

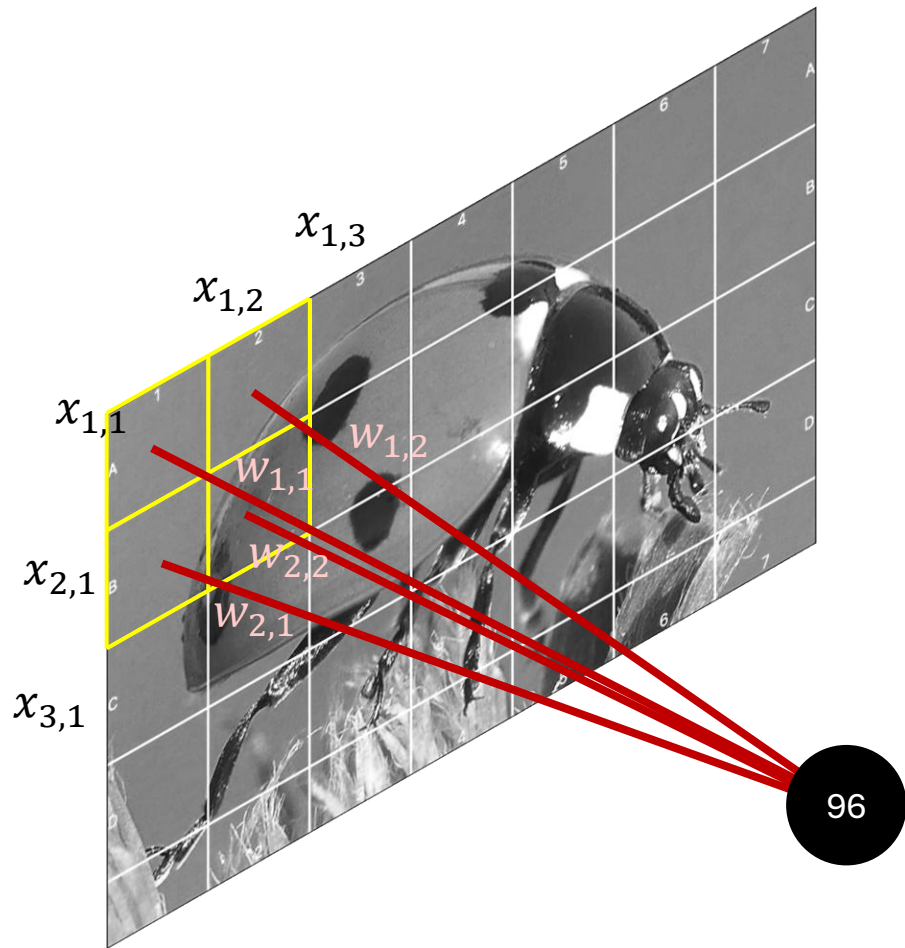
Muitos parâmetros =
Maior chance de *overfitting*
e alto custo computacional
(principalmente memória)

Translation Invariance




Localidade






$$y_{1,1} = \text{ReLU}(b + w_{1,1}x_{1,1} + w_{1,2}x_{1,2} + w_{2,1}x_{2,1} + w_{2,2}x_{2,2})$$

$$y_{1,1} = \text{ReLU}\left(b + \sum \sum x_{i,j} w_{i,j}\right)$$




$$y_{1,1} = \text{ReLU}(b + w_{1,1}x_{1,1} + w_{1,2}x_{1,2} + w_{2,1}x_{2,1} + w_{2,2}x_{2,2})$$

$$y_{1,1} = \text{ReLU}\left(b + \sum \sum x_{i,j}w_{i,j}\right)$$



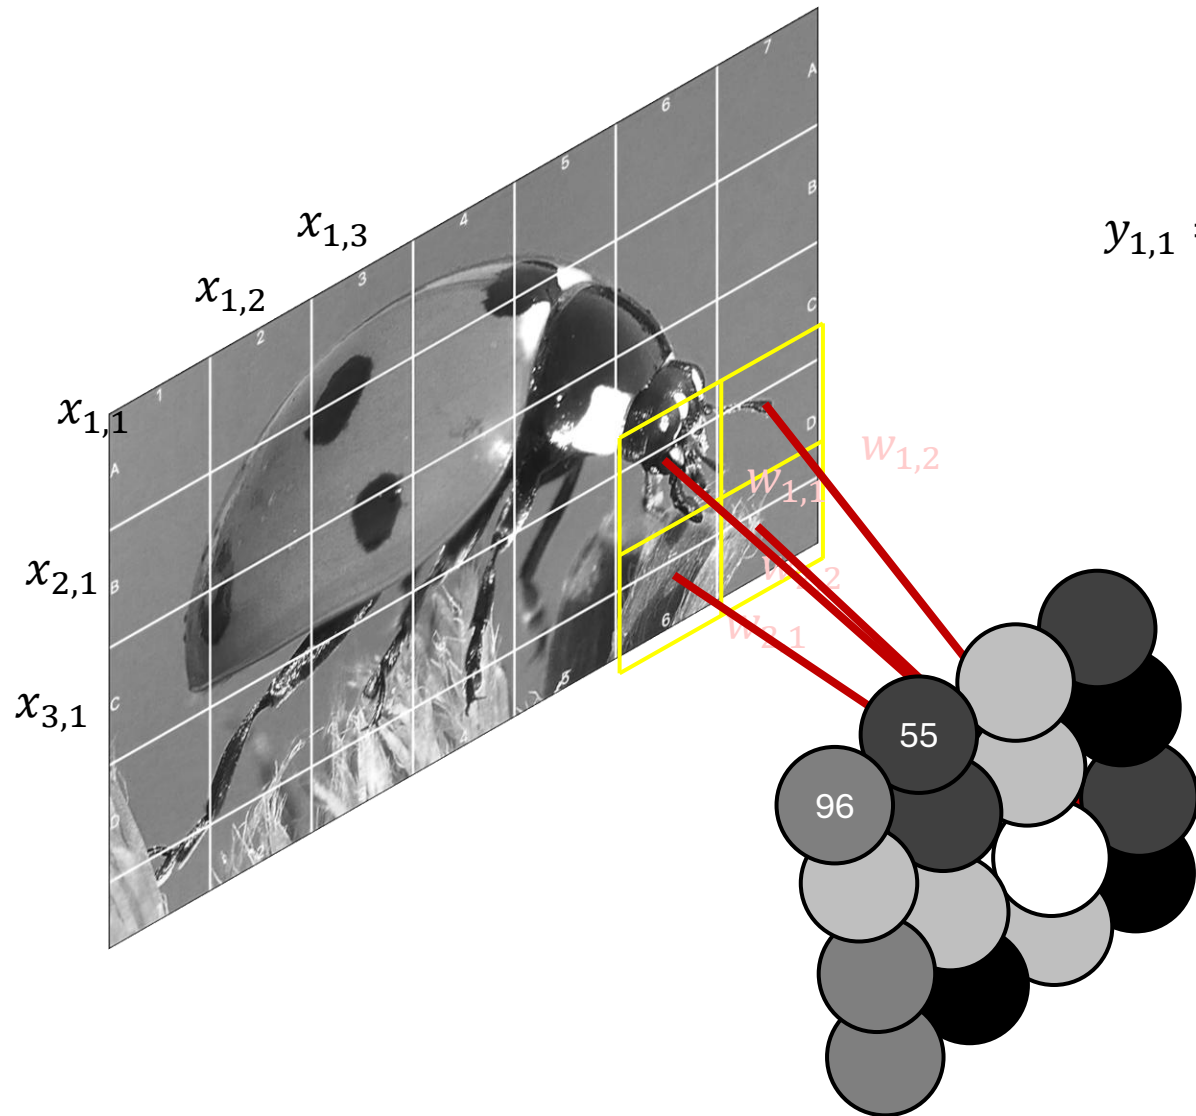
$$y_{1,1} = \text{ReLU}(b + w_{1,1}x_{1,1} + w_{1,2}x_{1,2} + w_{2,1}x_{2,1} + w_{2,2}x_{2,2})$$

$$y_{1,1} = \text{ReLU}\left(b + \sum \sum x_{i,j}w_{i,j}\right)$$





$$y_{1,1} = \text{ReLU}(b + w_{1,1}x_{1,1} + w_{1,2}x_{1,2} + w_{2,1}x_{2,1} + w_{2,2}x_{2,2})$$

$$y_{1,1} = \text{ReLU}\left(b + \sum \sum x_{i,j}w_{i,j}\right)$$



$$y_{1,1} = \text{ReLU}(b + w_{1,1}x_{1,1} + w_{1,2}x_{1,2} + w_{2,1}x_{2,1} + w_{2,2}x_{2,2})$$


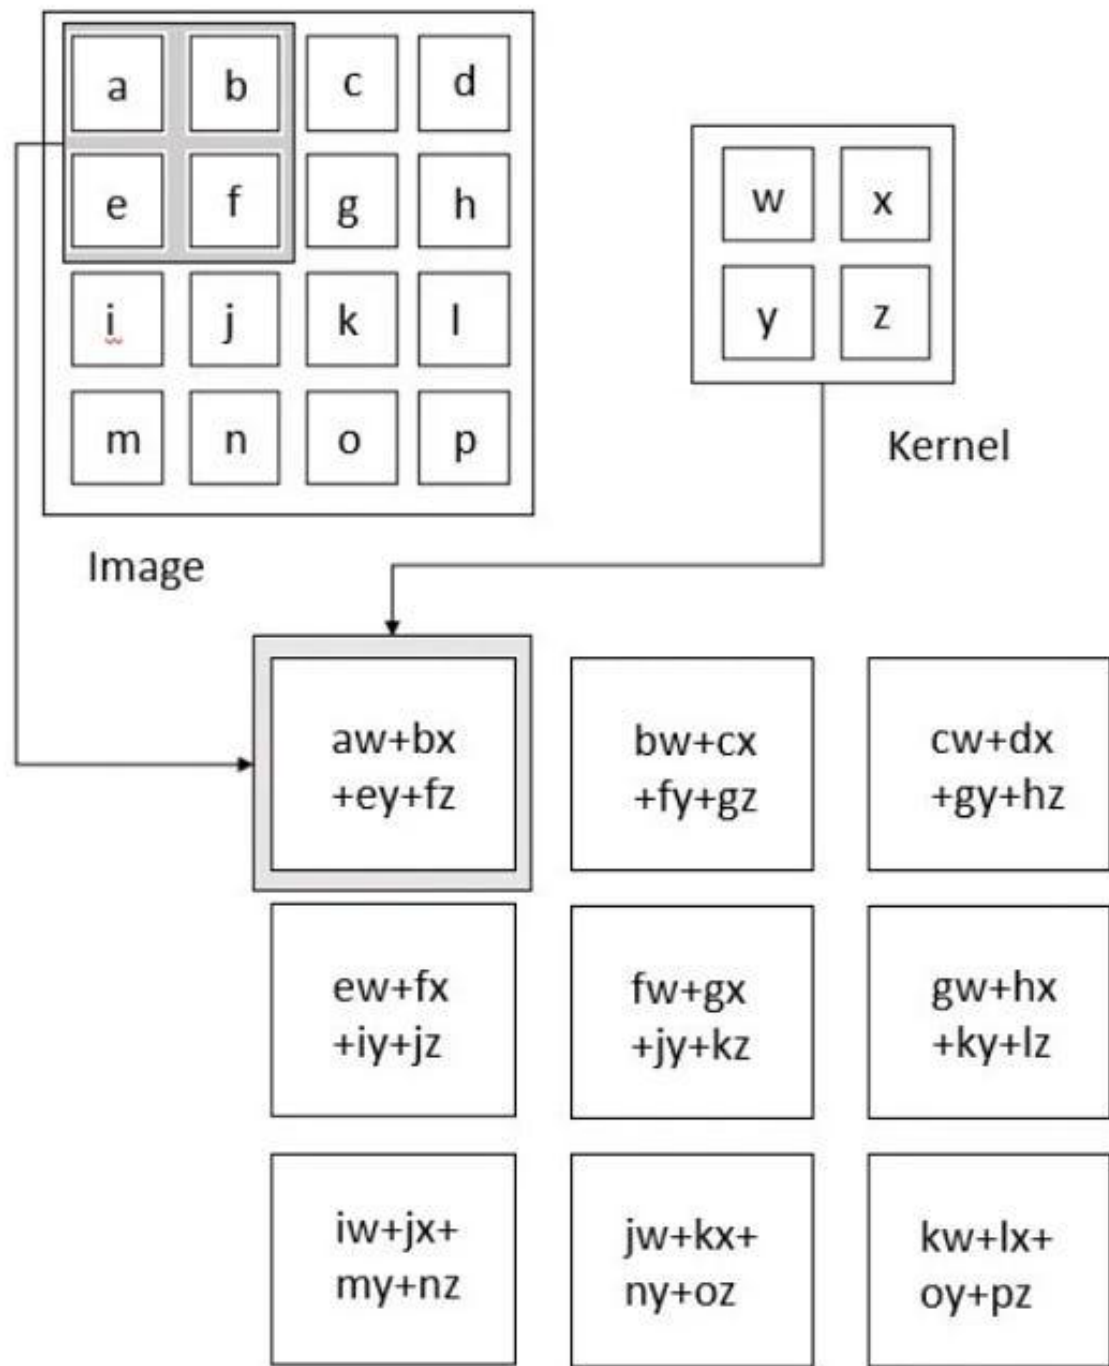
$$y_{1,1} = \text{ReLU}\left(b + \sum \sum x_{i,j}w_{i,j}\right)$$



Filtro ou kernel

$$y_{1,1} = \text{ReLU}(b + w_{1,1}x_{1,1} + w_{1,2}x_{1,2} + w_{2,1}x_{2,1} + w_{2,2}x_{2,2})$$

$$y_{1,1} = \text{ReLU}\left(b + \sum \sum x_{i,j} w_{i,j}\right)$$



Activation Map

Input Slice (6 x 6)

Receptive
Field

20	24	11	12	16	19
19	17	20	23	15	9
21	40	25	13	14	8
9	18	8	6	11	22
31	3	7	9	17	23
20	12	3	11	19	30

*

Sobel Kernel
(3 x 3)

1	0	-1
2	0	-2
1	0	-1

=

Output (4 x 4)

3			

Sample calculation for first filter location

$$\begin{aligned}
 &(20 \times 1) + (24 \times 0) + (11 \times -1) + \\
 &(19 \times 2) + (17 \times 0) + (20 \times -2) + \\
 &(21 \times 1) + (40 \times 0) + (25 \times -1) = 3
 \end{aligned}$$

20	24	11	12	16	19
19	17	20	23	15	9
21	40	25	13	14	8
9	18	8	6	11	22
31	3	7	9	17	23
20	12	3	11	19	30

*

1	0	-1
2	0	-2
1	0	-1

=

3	27		

Sample calculation for second filter location

$$\begin{aligned}
 &(24 \times 1) + (11 \times 0) + (12 \times -1) + \\
 &(17 \times 2) + (20 \times 0) + (23 \times -2) + \\
 &(40 \times 1) + (25 \times 0) + (13 \times -1) = 27
 \end{aligned}$$

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image


4		

Convolved
Feature

Input image: 5x5px




Filter




2x2



Feature map



4x4



3x3



3x3

Padding

0	0	0	0	0	0	0	0
0	3	3	4	4	7	0	0
0	9	7	6	5	8	2	0
0	6	5	5	6	9	2	0
0	7	1	3	2	7	8	0
0	0	3	7	1	8	3	0
0	4	0	4	3	2	2	0
0	0	0	0	0	0	0	0

$6 \times 6 \rightarrow 8 \times 8$

*


1	0	-1
1	0	-1
1	0	-1

3×3


=

-10	-13	1			
-9	3	0			

6×6




No Padding, Stride = 1



Same Padding, Stride = 1

Ver código de convolução no notebook



0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...


Input Channel #2 (Green)

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1




308

+

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2




-498

+

0	1	1
0	1	0
1	-1	1

Kernel Channel #3



164

+ 1 = -25



Bias = 1


Output

-25				...
				...
				...
				...
...

Input Tensor
 $w \times h \times 3$

Two $3 \times 3 \times 3$ Filters

Activation Map
 $w \times h \times 2$



Two filters (F_1 and F_2) contain three Kernels each for a total of 54 weights ($2 \times 3 \times 3 = 54$), plus a bias term for each filter = 56 trainable parameters

The output activation map now has a depth of two because two filters were convolved with the input.

Convolution Operation
for a Single Filter Location

Apply Activation

Kernel size, Padding e Stride

Veja [animação](#).


Pooling

1	3	1	5
4	0	3	3
6	1	5	1
1	4	0	8

Average Pooling



2	3
3	3.5




4	5
6	8


Max Pooling

Color Input Image
224 x 224 x 3

One 3 x 3 x 3 Filter




k_1 k_2 k_3




F_1

Activation
Function

$$f(wx + b)$$



Activation Map
224 x 224 x 1



Max
Pooling

112 x 112 x 1



This becomes the
input for the next
convolutional block

Downsized activation map
from Max Pooling operation

R G B

Convolution Operation

Apply Activation

Max Pooling

Convolutional Layer

Convolutional Block


Single filter (F_1) contains three Kernels (k_1, k_2, k_3)
for a total of 27 weights ($3 \times 3 \times 3 = 27$) plus
one bias term = 28 trainable parameters

Data Augmentation

Ajuda a prevenir *overfitting* e torna o modelo robusto à alguns tipos de ruído criando dados artificiais




Arquitectura Completa



Ver operações visualmente na [CNNExplainer](#)


Ver visualização em 3D

Ver como treinar uma rede com pytorch no notebook.




Input Image


Filters learn to detect structural patterns



F₁ = Horizontal Lines
F₂ = Vertical Lines
F₃ = Diagonal Lines



Activation Maps


 = Highly activated neurons in activation maps

Dê uma olhada no artigo da [AlexNet](#)

- Visualização de kernels aprendidos.
- *Embeddings*, distâncias entre *embeddings* e aprendizado de representações.
- Dropout, `net.train()` e `net.eval()`


Redes Neurais Pré-Treinadas, Transfer Learning, Feature Extraction, Fine-Tuning

Feature Extraction



In feature extraction, you **freeze the pre-trained model** layers to preserve existing learning and **add new layers** to learn additional information.

Fine Tuning



In fine-tuning, you **unfreeze the entire model** and **train it with a lower learning rate** to adapt to new challenges.

Ver notebook exemplo completo